

# The Sieve of Eratosthenes using MPI

Devin Delfino

Comp 401: Senior Seminar

4/16/2015

# Outline

1. Parallel Programming and MPI
2. The Sequential Algorithm
3. The Parallel Algorithm
4. Sequential vs. Parallel Comparison

# Outline

1. Parallel Programming and MPI
2. The Sequential Algorithm
3. The Parallel Algorithm
4. Sequential vs. Parallel Comparison

# Introduction

- **Parallel Computing** is the use of multiple computers or processors to reduce the time needed to solve a single computational problem.
- A **task** is a single program including local memory and a collection of input/output ports.
- A **channel** is a message queue between two tasks used for communication

# Ian Foster's Design Methodology

1. **Partitioning** - the process of dividing the computations and data into pieces.
2. **Communication** - channels between tasks allow communication between them
  - Local - a task's computation requires values from a small number of other tasks
  - Global - many tasks must contribute values to perform a computation
3. **Agglomeration** - grouping tasks in order to improve performance and reduce overhead.
4. **Mapping** - assigning processes or tasks to specific processors or computers

# Message Passing Interface (MPI)

- The most popular message-passing library standard
- There are many free implementations of MPI libraries, including OpenMPI and MPICH
- Integrates sequential language with functions that allow processes to communicate with each other

# Outline

1. Parallel Programming and MPI
2. The Sequential Algorithm
3. The Parallel Algorithm
4. Sequential vs. Parallel Comparison

# The Sequential Algorithm

---

## The Sieve of Eratosthenes

---

**Summary:** Finds all primes between 2 and  $n$ , inclusive

Create a list of natural numbers 2, 3, ... ,  $n$ , none of which are marked

Set  $k$  equal to the first prime number, 2

**while**  $k^2 \leq n$  **do**

    Mark all multiples of  $k$  between  $k^2$  and  $n$

    Set  $k$  to the smallest unmarked number greater than the current  $k$

**end while**

---



# Outline

1. Parallel Programming and MPI
2. The Sequential Algorithm
3. The Parallel Algorithm
4. Sequential vs. Parallel Comparison

# The Parallel Algorithm: Block Decomposition

- $n$  is the size of array and  $p$  is the number of processors
- Divide the array into  $p$  contiguous blocks of roughly equal size
- If  $n$  is divisible by  $p$ , then there will be  $p$  blocks of size  $n/p$ . Otherwise, there will be  $p$  blocks of with sizes of either  $\lfloor n/p \rfloor$  or  $\lceil n/p \rceil$ .
- Common data block computations include finding the first/last element controlled by a given process

# The Parallel Algorithm: Block Decomposition, cont.

- $n$  Suppose is the size of array and  $p$  is the number of processors
- First element controlled by process  $i$ :

$$\lfloor in/p \rfloor$$

- Last element controlled by process  $i$ :

$$\lfloor (i+1)n/p \rfloor - 1$$

# Block Decomposition Example, $n = 121$ , $p = 3$

## Process 0:

- First index:  $\lfloor (0)(121)/3 \rfloor = 0$
- Last index:  $\lfloor (0 + 1)121/3 \rfloor - 1 = \lfloor 40\frac{1}{3} \rfloor - 1 = 39$
- Size: last - first + 1 =  $\lfloor n/p \rfloor = 40$ .

## Process 1:

- First index:  $\lfloor (1)(121)/3 \rfloor = \lfloor 40\frac{1}{3} \rfloor = 40$
- Last index:  $\lfloor (1 + 1)121/3 \rfloor - 1 = \lfloor 80\frac{2}{3} \rfloor - 1 = 79$
- Size: last - first + 1 = 40 =  $\lfloor n/p \rfloor$ .

## Process 2:

- First index:  $\lfloor (2)(121)/3 \rfloor = \lfloor 80\frac{2}{3} \rfloor = 80$
- Last index:  $\lfloor (2 + 1)121/3 \rfloor - 1 = \lfloor 121 \rfloor - 1 = 120$
- Size: last - first + 1 = 41 =  $\lceil n/p \rceil$ .

# Developing the Algorithm, Step 1

1. Create a list of natural numbers 2, 3, ... ,  $n$ , none of which are marked
  - Each task handles a specific block of the entire array
  - Use the formulas to determine the numbers represented by the first and last elements of the block along with its size
  - Keep in mind the difference between the local index for the block and the global index for the entire array
  - To minimize communication between tasks, make task 0 responsible for finding the next value of  $k$
  - This can be done by ensuring  $n/p > \sqrt{n}$

## 2. Set $k$ equal to the first prime number, 2

- $k$  is set to the first prime number for all tasks
- After each iteration, each task must be told the next value of  $k$

## Developing the Algorithm, Step 3

While  $k^2 \leq n$

3a. Mark all multiples of  $k$  between  $k^2$  and  $n$

- Must determine the first multiple of  $k$  in the given block (if it's greater than  $k^2$ )
- From the first multiple of  $k$ , mark every  $k^{th}$  element

3b. Set  $k$  to the smallest unmarked number greater than the current  $k$

- The smallest unmarked number greater than the current  $k$  is always part of the block belonging to task 0
- Task 0 finds the next value of  $k$ , and broadcasts it to the rest of the tasks

# Outline

1. Parallel Programming and MPI
2. The Sequential Algorithm
3. The Parallel Algorithm
4. Sequential vs. Parallel Comparison



# Sequential vs. Parallel

# References

- ▶ MICHAEL J. QUINN.  
*Parallel Programming in C with MPI and OpenMP.*  
McGraw-Hill Science/Engineering/Math, 2003.